COMP90020 Distributed Algorithms

Project Demo: Dynamic Vector Clocks

Team Double-J
James Sammut (502030)
Joel Kenna (995401)

Introduction and Background

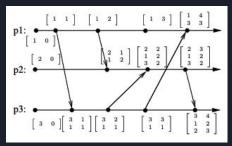
- The concept of "Logical Time": mapping the physical phenomenon when utilised with distributed systems
- Logical clocks the ordering of events within processes part of a distributed system. Lamport & Vector Clocks [3]-[5] are famous and well-regarded in this context
- Various extensive clock-based algorithms exist:, Dynamic
 Vector Clocks [1], Interval Tree Clocks [6], Prime Clocks [7],
 Chain Clocks [8], Matrix Clocks [1] [9] [10]
- We've chosen to implement Dynamic Vector Clocks for this project!



Clocks, clocks and clocks!

Dynamic Vector Clocks

- Vector Clocks: an array (V) of scalar values.
 - V[i] Incremented by P_i upon message send/internal event
 - Received by process P_i who increments V[j]
- The number of processes (n) is pre-determined: not adjustable!
- Dynamic Vector Clocks (DVCs) are a way of making this requirement unnecessary: processes (peers) join a session and combine their vector clocks on message delivery.
- Peers (P_i) hold onto a clock with $[P_i, <int>]$ for all $P_j: P_i \neq P_j$ for each other process P_i it knows of after message receive/delivery.
- The DVC algorithms allows flexibility and is the basis of our projects implementation



Dynamic Vector Clock Process Diagram [2]

Dynamic Vector Clocks - The Rationale

Causal Broadcast Ordering

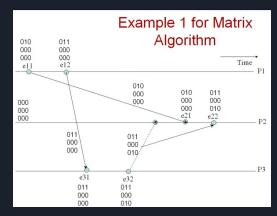
- Non-causal ordering of messages = **confusion**.
- A process lags/network becomes segregated unordered.

Dynamic Networks and Peer Size

- For modern P2P Messaging Applications peer size not fixed.
- Peers should be able to join/leave whenever they like.
- A DVC can grow/shrink as needed.

Simpler Implementation

- Peers only need to hold onto a clock with [Process_ID, clock]
- Matrix Clocks initially considered.
 - \circ **Huge overhead** with storing sends for each send to P_i from P_i
 - o Grows exponentially (i.e 20 peers, 400 values)
 - Not applicable for our use case of **broadcast-only**
 - Delivery logic is similar to DVCs.



Matrix Clock and Broadcast between Processes [1]

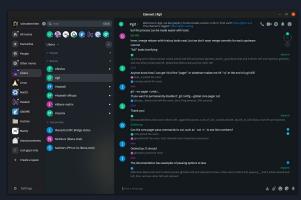
Our Application - P2P Messaging Application

Application Features

- P2P messaging application for scenarios with low room for confusion
- Peers are able to broadcast messages to one-another
- Guaranteed causal message delivery
- Peers can **enter** and **leave** the network dynamically
- Simple GUI (Kivy library) for improved user experience

Implemented in 2 Phases

- **Phase 1**: Algorithm development using MPI
 - Dynamic Vector Clock & Matrix Clock algorithm implementations
- Phase 2: Algorithm socket integration and GUI
 - Utilising algorithm implemented from Phase #1
 - Allowance for peers to enter/leave as needed
 - With a peer registry server, or without



Chat Room - just an example!

Phase 1: Algorithm development using MPI

- Implemented using Python & MPI (Message Passing Interface)
- Dynamic Vector Clock and Matrix Clock algorithms
- Simulated environment of *n* processes (whom are 'aware' of each other)
- Both broadcast and unicast messages.
- Input of events on a process (text file) -> output on terminal with causal delivery checking.
- Enqueued if deliverability is not adhered.

```
1 b1, r3, r2
2 r1, b2, b3
3 r2, r3, r1
```

Process events - invalid broadcast example

```
./phase1_invoke.sh -f examples/dynamic_vector_clocks/b1_3_node_simple_valid.txt -a dvc
Event #0 -> b1: (0)
Broadcast message to process(es) [2, 3]
Process 1 generated message heading for Process(es) [2, 3], Number: 9.936, DVC of [[1, 1], [2, 0], [3, 0]]
Process 1 broadcasting message to Process(es) [2, 3] @ 22:02:44.117457
Event #0 -> r1: (0)
Process 2 received number 9.936 from Process 1 @ 22:02:44.117602
This message satisfied the DVC causal deliverability condition. Delivering.
3.936 (message number) + 0 (current sum). Process 2's sum = 9.936
Checking messages in the message/hold back queue for deliverability
No messages are in the message/hold back gueue
Event #0 -> r1: (0)
Process 3 received number 9.936 from Process 1 @ 22:02:44.117617
This message satisfied the DVC causal deliverability condition. Delivering.
3.936 (message number) + 0 (current sum), Process 3's sum = 9.936
Checking messages in the message/hold back queue for deliverability
No messages are in the message/hold back queue
Number Sum: 9.936
```

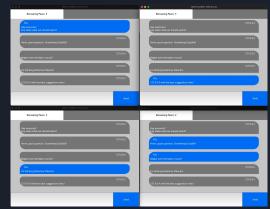
Phase 1 invocation - valid broadcast example

Phase 2: Algorithm socket integration & GUI

- Implemented using Python & sockets
- Dynamic Vector Clock algorithm implemented from Phase 1 integrated
- Processes ("peers") able to dynamically join/leave the system
- All processes able to communicate with one-another: a "chat room"
- Implementation with a peer registry server (for handling registration), or not (specify IPs for peers)
- Messages that are broadcast: checked for causal delivery. Enqueued if deliverability is not adhered.



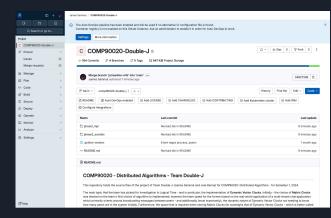
Phase 2 invocation - without peer registry server



Phase 2 invocation - with peer registry server

COMP90020-Double-J Repository

- This project's implementation of both phases has been implemented within a <u>repository</u> on the Melbourne School of Engineerings' GitLab instance.
- phase1_mpi/and phase2_sockets/directories implementations for each phase separated
- The README . md within the repository contains detailed information for both phases:
 - Approach and background
 - Implementation specifics
 - Invocation
- Our code submission includes all repository files!



COMP90020-Double-J Repository



Thank you!

References

- [1] T. Landes. (2006). Dynamic Vector Clocks for Consistent Ordering of Events in Dynamic Distributed Applications. Presented at Proceedings of the International Conference on Parallel and Distributed Processing Techniques and Applications & Conference on Real-Time Computing Systems and Applications [Online]. Available: https://vs.inf.ethz.ch/edu/HS2017/VS/exercises/A3/DVC_Landes.pdf
- [2] N. Meghanathan. Module 6.2.3 Matrix Algorithm Causal Delivery of Messages. (Nov. 12, 2013). Accessed: Mar. 13, 2024. [Online video]. Available: https://www.youtube.com/watch?v=WgTx7BHWzts.
- [3] L. Lamport, "Time, clocks, and the ordering of events in a distributed system," in *Communications of the ACM*, vol.21, no. 7, Jul. 1, 1978, pp. 558-565, doi: 10.1145/359545.359563.
- [4] P. Redmond, G. Shen, N. Vazou and L. Kuper, "Verified Causal Broadcast with Liquid Haskell," in *IFL* '22: *Proceedings of the 34th Symposium on Implementation and Application of Functional Languages*, Copenhagen, Denmark, 2022, pp. 1-13, doi: 10.1145/3587216.3587222.
- [5] M. Raynal and M. Singhal, "Logical time: capturing causality in distributed systems," in *Computer*, Feb. 1996, vol. 29, no. 2, pp. 49-56,, doi: 10.1109/2.485846.
- [6] P. Sérgio Almeida, C. Baquero and V Fonte, "Interval Tree Clocks," in *OPODIS '08: Proceedings of the 12th International Conference on Principles of Distributed Systems*, Luxor, Egypt, Dec. 2008, pp. 259-274. doi: 10.1007/978-3-540-92221-6_18.
- [7] A. D. Kshemkalyani, M Shen and B. Voleti, "Prime clock: Encoded vector clock to characterize causality in distributed systems" in *Journal of Parallel and Distributed Computing*, vol. 140, pp. 37-51, doi: 10.1016/j.jpdc.2020.02.008.
- [8] A. Agarwal and V. K. Garg, "Efficient Dependency Tracking for Relevant Events in Shared-Memory Systems" in PODC '05: Proceedings of the twenty-fourth annual ACM symposium on Principles of distributed computing, New York, NY, USA, Jul 17, 2005, pp. 19-28, doi: 10.1145/1073814.1073818.
- [9] A. Singh and N. Badal, "An efficient implementation of multi matrix clocks in the distributed system," in 2015 International Conference on Green Computing and Internet of Things (ICGCIoT), Greater Noida, India, Oct. 8-10, 2015, pp. 209-213, doi: 10.1109/ICGCIoT.2015.7380459.
- [10]] H. Guerreiro, L. Rodrigues, N. Preguiça and N. Quental, "Causality Tracking Trade-offs for Distributed Storage," in 2020 IEEE 19th International Symposium on Network Computing and Applications (NCA), Cambridge, MA, USA, Jun. 24-27, 2020, pp. 1-10, doi: 10.1109/NCA51143.2020.9306734.