



COMP90020

Distributed Algorithms

Project Demo: Dynamic Vector
Clocks

Team Double-J
James Sammut (502030)
Joel Kenna (995401)

Introduction and Background


- The concept of “Logical Time”: mapping the physical phenomenon when utilised with distributed systems
- Logical clocks - the ordering of events within processes part of a distributed system. **Lamport & Vector Clocks** [3]-[5] are famous and well-regarded in this context
- Various extensive clock-based algorithms exist:, **Dynamic Vector Clocks** [1], **Interval Tree Clocks** [6], **Prime Clocks** [7], **Chain Clocks** [8], **Matrix Clocks** [1] [9] [10]
- We've chosen to implement **Dynamic Vector Clocks** for this project!



Clocks, clocks and clocks!

Dynamic Vector Clocks

- Vector Clocks: an array (V) of scalar values.
 - $V[i]$ Incremented by P_i upon message send/internal event
 - Received by process P_j , who increments $V[j]$
- **The number of processes (n) is pre-determined: not adjustable!**
- Dynamic Vector Clocks (DVCs) are a way of making this requirement unnecessary: processes (peers) join a session and combine their vector clocks on message delivery.
- Peers (P_i) hold onto a clock with $[P_i, <int>]$ for all $P_j: P_i \neq P_j$ for each other process P_j it knows of after message receive/delivery.
- The DVC algorithms allows flexibility and is the basis of our projects implementation



Dynamic Vector Clock Process Diagram [2]

Dynamic Vector Clocks - The Rationale

Causal Broadcast Ordering

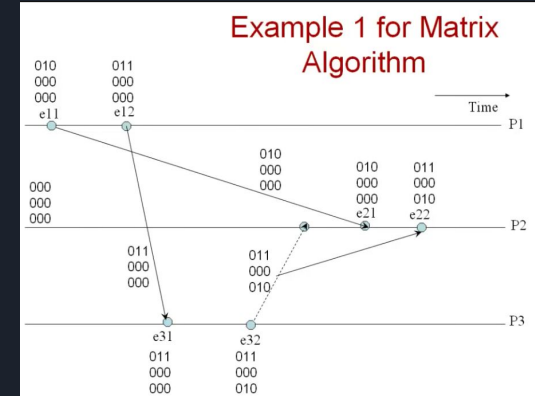
- Non-causal ordering of messages = **confusion**.
- A process lags/network becomes segregated - **unordered**.

Dynamic Networks and Peer Size

- For modern P2P Messaging Applications - peer size not fixed.
- Peers **should be able to join/leave whenever they like**.
- A DVC can grow/shrink as needed.

Simpler Implementation

- Peers only need to hold onto a clock with [Process_ID, clock]
- Matrix Clocks - **initially considered**.
 - **Huge overhead** with storing sends for each send to P_j from P_i
 - Grows exponentially (i.e 20 peers, 400 values)
 - Not applicable for our use case of **broadcast-only**
 - Delivery logic is similar to DVCs.



Matrix Clock and Broadcast between Processes [1]

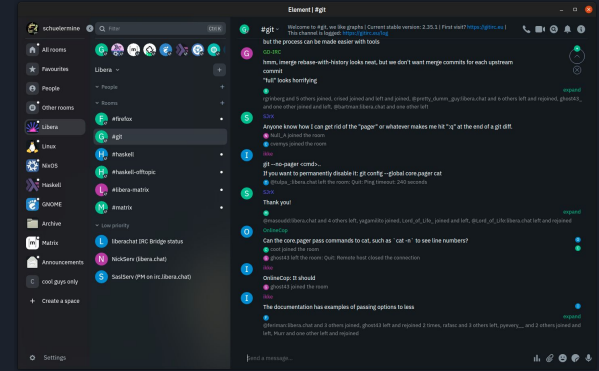
Our Application - P2P Messaging Application

Application Features

- P2P messaging application for scenarios with low room for confusion
- Peers are able to broadcast messages to one-another
- **Guaranteed** causal message delivery
- Peers can **enter** and **leave** the network dynamically
- Simple GUI (Kivy library) for improved user experience

Implemented in 2 Phases

- **Phase 1:** Algorithm development using MPI
 - Dynamic Vector Clock & Matrix Clock algorithm implementations
- **Phase 2:** Algorithm socket integration and GUI
 - Utilising algorithm implemented from Phase #1
 - Allowance for peers to enter/leave as needed
 - With a peer registry server, or without



Chat Room - just an example!

Phase 1: Algorithm development using MPI

- Implemented using Python & MPI (Message Passing Interface)
- Dynamic Vector Clock and Matrix Clock algorithms
- Simulated environment of n processes (whom are 'aware' of each other)
- Both broadcast and unicast messages.
- Input of events on a process (text file) -> output on terminal with causal delivery checking.
- Enqueued if deliverability is not adhered.

1 b1, r3, r2
2 r1, b2, b3
3 r2, r3, r1

Process events - invalid broadcast example

```
./phase1_invoke.sh -f examples/dynamic_vector_clocks/bl_3_node_simple_valid.txt -a dvc
|----- Process 1: ['b1'] -----|
Event #0 -> b1: (0)
Broadcast message to process(es) [2, 3]
Process 1 generated message heading for Process(es) [2, 3]. Number: 9.936. DVC of [[1, 1], [2, 0], [3, 0]]
Process 1 broadcasting message to Process(es) [2, 3] @ 22:02:44.117457

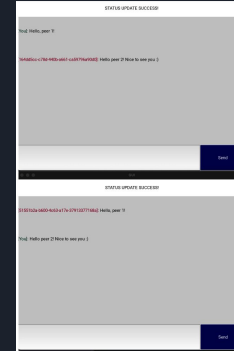
|----- Process 2: ['r1'] -----|
Event #0 -> r1: (0)
Process 2 received number 9.936 from Process 1 @ 22:02:44.117602
This message satisfied the DVC causal deliverability condition. Delivering.
9.936 (message number) + 0 (current sum). Process 2's sum = 9.936
Checking messages in the message/hold back queue for deliverability
No messages are in the message/hold back queue
WC after r1: [[1, 1], [2, 0], [3, 0]]
Number Sum: 9.936

|----- Process 3: ['r1'] -----|
Event #0 -> r1: (0)
Process 3 received number 9.936 from Process 1 @ 22:02:44.117617
This message satisfied the DVC causal deliverability condition. Delivering.
9.936 (message number) + 0 (current sum). Process 3's sum = 9.936
Checking messages in the message/hold back queue for deliverability
No messages are in the message/hold back queue
WC after r1: [[1, 1], [2, 0], [3, 0]]
Number Sum: 9.936
```

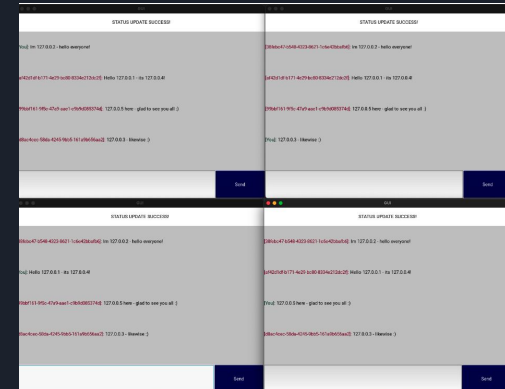
Phase 1 invocation - valid broadcast example

Phase 2: Algorithm socket integration & GUI

- Implemented using Python & sockets
- Dynamic Vector Clock algorithm implemented from Phase 1 integrated
- Processes (“peers”) able to dynamically join/leave the system
- All processes able to communicate with one-another: a “chat room”
- Implementation with a **peer registry server** (for handling registration), or not (specify IPs for peers)
- Messages that are broadcast: checked for **causal delivery**. Enqueued if deliverability is not adhered.



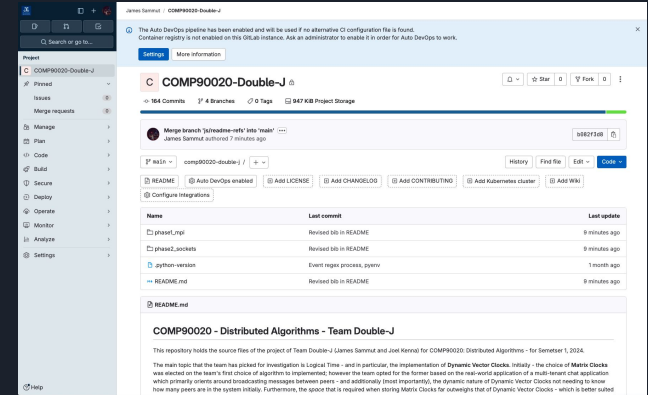
Phase 2 invocation - without peer registry server



Phase 2 invocation - with peer registry server

COMP90020-Double-J Repository

- This project's implementation of both phases has been implemented within a [repository](#) on the Melbourne School of Engineerings' GitLab instance.
- `phase1_mpi/` and `phase2_sockets/directories` - implementations for each phase separated
- The `README.md` within the repository contains detailed information for both phases:
 - Approach and background
 - Implementation specifics
 - Invocation
- Our code submission includes all repository files!



COMP90020-Double-J Repository



Demo Time





Thank you!



References

- [1] T. Landes. (2006). Dynamic Vector Clocks for Consistent Ordering of Events in Dynamic Distributed Applications. Presented at Proceedings of the International Conference on Parallel and Distributed Processing Techniques and Applications & Conference on Real-Time Computing Systems and Applications [Online]. Available: https://vs.inf.ethz.ch/edu/HS2017/VS/exercises/A3/DVC_Landes.pdf
- [2] N. Meghanathan. Module 6.2.3 Matrix Algorithm Causal Delivery of Messages. (Nov. 12, 2013). Accessed: Mar. 13, 2024. [Online video]. Available: <https://www.youtube.com/watch?v=WgTx7BHWzts>.
- [3] L. Lamport, "Time, clocks, and the ordering of events in a distributed system," in *Communications of the ACM*, vol.21, no. 7, Jul. 1, 1978, pp. 558-565, doi: 10.1145/359545.359563.
- [4] P. Redmond, G. Shen, N. Vazou and L. Kuper, "Verified Causal Broadcast with Liquid Haskell," in *IFL '22: Proceedings of the 34th Symposium on Implementation and Application of Functional Languages*, Copenhagen, Denmark, 2022, pp. 1-13, doi: 10.1145/3587216.3587222.
- [5] M. Raynal and M. Singhal, "Logical time: capturing causality in distributed systems," in *Computer*, Feb. 1996, vol. 29, no. 2, pp. 49-56., doi: 10.1109/2.485846.
- [6] P. Sérgio Almeida, C. Baquero and V. Fonte, "Interval Tree Clocks," in *OPODIS '08: Proceedings of the 12th International Conference on Principles of Distributed Systems*, Luxor, Egypt, Dec. 2008, pp. 259-274. doi: 10.1007/978-3-540-92221-6_18.
- [7] A. D. Kshemkalyani, M. Shen and B. Voleti, "Prime clock: Encoded vector clock to characterize causality in distributed systems" in *Journal of Parallel and Distributed Computing*, vol. 140, pp. 37-51, doi: 10.1016/j.jpdc.2020.02.008.
- [8] A. Agarwal and V. K. Garg, "Efficient Dependency Tracking for Relevant Events in Shared-Memory Systems" in *PODC '05: Proceedings of the twenty-fourth annual ACM symposium on Principles of distributed computing*, New York, NY, USA, Jul 17, 2005, pp. 19-28, doi: 10.1145/1073814.1073818.
- [9] A. Singh and N. Badal, "An efficient implementation of multi matrix clocks in the distributed system," in *2015 International Conference on Green Computing and Internet of Things (ICGCIoT)*, Greater Noida, India, Oct. 8-10, 2015, pp. 209-213, doi: 10.1109/ICGCIoT.2015.7380459.
- [10] H. Guerreiro, L. Rodrigues, N. Preguiça and N. Quental, "Causality Tracking Trade-offs for Distributed Storage," in *2020 IEEE 19th International Symposium on Network Computing and Applications (NCA)*, Cambridge, MA, USA, Jun. 24-27, 2020, pp. 1-10, doi: 10.1109/NCA51143.2020.9306734.